

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 01

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
data=pd.read_csv('/content/Iris_Dataset.csv')
data
```

```
   Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm variety
0  1  5.1  3.5  1.4  0.2  Iris-setosa
1  2  4.9  3.0  1.4  0.2  Iris-setosa
2  3  4.7  3.2  1.3  0.2  Iris-setosa
3  4  4.6  3.1  1.5  0.2  Iris-setosa
4  5  5.0  3.6  1.4  0.2  Iris-setosa
... ..
145 146 6.7  3.0  5.2  2.3  Iris-virginica
146 147 6.3  2.5  5.0  1.9  Iris-virginica
147 148 6.5  3.0  5.2  2.0  Iris-virginica
148 149 6.2  3.4  5.4  2.3  Iris-virginica
149 150 5.9  3.0  5.1  1.8  Iris-virginica
150 rows × 6 columns
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 # Column Non-Null Count Dtype
---  ---
0    Id 150 non-null int64
1    SepalLengthCm 150 non-null float64
2    SepalWidthCm 150 non-null float64
3    PetalLengthCm 150 non-null float64
```

```

4   PetalWidthCm 150 non-null float64
5   variety      150 non-null object   dtypes:
float64(4), int64(1), object(1)  memory usage:
7.2+ KB  data.describe()

```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	count
	150.000000	150.000000	150.000000	150.000000	150.000000	
mean	75.500000	5.843333	3.054000	3.758667	1.198667	std
	43.445368	0.828066	0.433594	1.764420	0.763161	
min	1.000000	4.300000	2.000000	1.000000	0.100000	
25%	38.250000	5.100000	2.800000	1.600000	0.300000	
50%	75.500000	5.800000	3.000000	4.350000	1.300000	
75%	112.750000	6.400000	3.300000	5.100000	1.800000	
max	150.000000	7.900000	4.400000	6.900000	2.500000	


```
data.value_counts('variety')
```

```

count
variety
Iris-setosa 50
Iris-versicolor 50
Iris-virginica 50

```

```
sns.countplot(x='variety',data=data,) plt.show()
```



```

dummies=pd.get_dummies(data.variety)
FinalDataset=pd.concat([pd.get_dummies(data.variety),data.iloc[:,[0,1,2,3]]], axis=1)

```

```
FinalDataset.head()
```

```


Iris-setosa Iris-versicolor Iris-virginica Id
SepalLengthCm SepalWidthCm PetalLengthCm 0 True False
False 1 5.1 3.5 1.4 1 True False False 2 4.9 3.0 1.4 2 True False
False 3 4.7 3.2 1.3 3 True False False 4 4.6 3.1 1.5 4
True False False 5 5 0 3 6 1 4

```


```

sns.scatterplot(x='SepalLengthCm',y='SepalWidthCm',hue='variety',data=data,)
<Axes: xlabel='SepalLengthCm', ylabel='SepalWidthCm'>

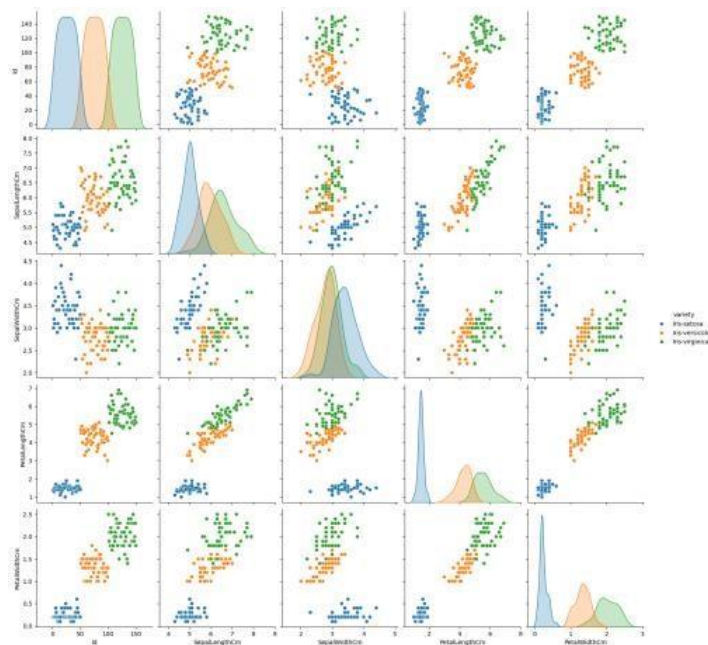
```




```
sns.scatterplot(x='PetalLengthCm',y='PetalWidthCm',hue='variety',data=data,)
```




```
sns.pairplot(data,hue='variety',height=3);
```




```
plt.show()
sns.FacetGrid(data,hue='variety',height=5).map(
sns.histplot,'Petal.LengthCm').add_legend();
plt.show();
```



```
sns.FacetGrid(data,hue='variety',height=5).map(
sns.histplot,'Petal.WidthCm').add_legend();
plt.show();
```



```
sns.FacetGrid(data,hue='variety',height=5).map(
sns.histplot,'SepalLengthCm').add_legend();
plt.show();
```



```
sns.FacetGrid(data,hue='variety',height=5).map(sns.histplot,'SepalWidthCm').a
dd_legend();
plt.show();
```

Lab experiments
Roll no: 230701038
Name: Arvind Ravi
Class: CSE-A III
Subject: Fundamentals of data science (CS2334)
Experiment: 02

```
import numpy as np
array=np.random.randint(1,100,9)  array  array([83, 25, 19,
47, 62, 15, 96, 39, 51])
np.sqrt(array)
array([9.11043358, 5. , 4.35889894, 6.8556546 , 7.87400787,
3.87298335, 9.79795897, 6.244998 , 7.14142843])
array.ndim
1      new_array=array.re
shape(3,3)  new_array
array([[83, 25, 19],
[47, 62, 15],
[96, 39, 51]])
new_array.ndim
2      new_array.ravel()
array([83, 25, 19, 47, 62, 15, 96, 39, 51])
newm=new_array.reshape(3,3)  newm
array([[83, 25, 19],
[47, 62, 15],
[96, 39, 51]])
newm[2,1:3]  array([39,
51])  newm[1:2,1:3]
array([[62, 15]])
new_array[0:3,0:0]  array([], shape=(3,
0), dtype=int64)  new_array[0:2,0:1]
```

```

    array([[83],[47]])
new_array[0:3,0:1]
    array([[83],[47],
    [96]])
new_array[1:3]
    array([[47, 62, 15],
    [96, 39, 51]])

```

Lab experiments
Roll no: 230701038
Name: Arvind Ravi
Class: CSE-A III
Subject: Fundamentals of data science (CS2334)
Experiment: 03

```
import numpy as np
```

```
import pandas as pd
list=[[1,'Smith',50000],[2,'Jones',60000]]
```

```
df=pd.DataFrame(list) df
```

```

    0 1 2
0 1 Smith 50000
1 2 Jones 60000

```

```
df.columns=['Empd','Name','Salary'] df
```

```

    Empd Name Salary
0 1 Smith 50000
1 2 Jones 60000

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2 entries, 0 to 1 Data
columns (total 3 columns):
# Column Non-Null Count Dtype
-----
0      Empd 2 non-null int64

```

```

1      Name 2 non-null object    2 Salary 2
non-null int64    dtypes: int64(2), object(1)
memory usage: 176.0+ bytes
df=pd.read_csv("/content/50_Startups.csv")
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49  Data
columns (total 5 columns):
# Column Non-Null Count Dtype
---  -
0    R&D Spend 50 non-null float64
1    Administration 50 non-null float64
2    Marketing Spend 50 non-null float64
3    State 50 non-null object    4 Profit 50
non-null float64  dtypes: float64(4),
object(1)  memory usage: 2.1+ KB  df.head()

```

```

      R&D Spend Administration Marketing Spend State Profit
0  165349.20 136897.80 471784.10 New York 192261.83
1  162597.70 151377.59 443898.53 California 191792.06
2  153441.51 101145.55 407934.54 Florida 191050.39
3  144372.41 118671.85 383199.62 New York 182901.99
4  142107.34 91391.77 366168.42 Florida 166187.94

```

```
df.tail()
```

```

      R&D Spend Administration Marketing Spend State Profit
45  1000.23 124153.04 1903.93 New York 64926.08
46  1315.46 115816.21 297114.46 Florida 49490.75
47  0.00 135426.92 0.00 California 42559.73
48  542.05 51743.15 0.00 New York 35673.41
49  0 00 116983.80 45173.06 California 14681.40

```

```

import numpy as np
import pandas as pd
df=pd.read_csv("/content/employee.csv")

```

```
df.head()
```

```
emp id name salary
```



```
0 1 SREE VARSSINI K S 5000
1 2 SREEMATHI B 6000
2 3 SREYA G 7000
3 4 SREYASKARI MULLAPUDI 5000
4 5 SRI AKASH U G 8000
```

```
df.tail()
   emp id name salary
2 3 SREYA G 7000
3 4 SREYASKARI MULLAPUDI 5000
4 5 SRI AKASH U G 8000
5 6 SRI HARSHAVARDHANAN R 3000
6 7 SRI HARSHAVARDHANAN R 6000
```

```
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7 entries, 0 to 6   Data
columns (total 3 columns):
# Column Non-Null Count Dtype
---  ---
0    emp id 7 non-null int64
1    name 7 non-null object    2
salary 7 non-null int64    dtypes:
int64(2), object(1)  memory usage:
296.0+ bytes  df.salary
```

```
   salary
0 5000
1 6000
2 7000
3 5000
4 8000
5 3000
6 6000
```

```
type(df.salary)
```

```
pandas.core.series.Series
```

```
def __init__(data=None, index=None, dtype: Dtype | None=None,
name=None, copy: bool | None=None, fastpath: bool=False) -> None
```

One-dimensional ndarray with axis labels (including time series).

Labels need not be unique but must be a hashable type. The object supports both integer- and label-based indexing and provides a host of methods for performing operations involving the index.

```
df.salary.mean()
```

```
5714.285714285715 df.salary.median()
```

```
6000.0 df.salary.mode()
```

```
salary
```

```
0 5000
```

```
1 6000
```

```
df.salary.var()
```

```
2571428.5714285714 df.salary.std()
```

```
1603.5674514745463 df.describe()
```

```
emp id salary
```

```
count 7.000000 7.000000
```

```
mean 4.000000 5714.285714
```

```
std 2.160247 1603.567451 min
```

```
1.000000 3000.000000
```

```
25% 2.500000 5000.000000
```

```
50% 4.000000 6000.000000
```

```
75% 5.500000 6500.000000 max
```

```
7 000000 8000 000000
```

```
df.describe(include='all')
```

```
emp id name salary
```

```
count 7.000000 7 7.000000
```

```
unique NaN 6 NaN top NaN SRI
```

```
HARSHAVARDHANAN R NaN
```

```
freq NaN 2 NaN
```

```
mean 4.000000 NaN 5714.285714
```

```
std 2.160247 NaN 1603.567451 min
```

```
1.000000 NaN 3000.000000 25%
```

```
2.500000 NaN 5000.000000
```

50% 4.000000 NaN 6000.000000
75% 5.500000 NaN 6500.000000
max 7 000000 NaN 8000 000000

```
empCol=df.columns empCol
```

```
Index(['emp id', 'name ', 'salary'], dtype='object')
```

```
emparray=df.values emparray array([[1, 'SREE VARSSINI  
K S', 5000],
```

```
[2, 'SREEMATHI B', 6000],  
[3, 'SREYA G', 7000],  
[4, 'SREYASKARI MULLAPUDI', 5000],  
[5, 'SRI AKASH U G', 8000],  
[6, 'SRI HARSHAVARDHANAN R', 3000],  
[7, 'SRI HARSHAVARDHANAN R', 6000]], dtype=object)
```

```
employee_DF=pd.DataFrame(emparray,columns=empCol) employee_DF
```

	emp id	name	salary
0	1	SREE VARSSINI K S	5000
1	2	SREEMATHI B	6000
2	3	SREYA G	7000
3	4	SREYASKARI MULLAPUDI	5000
4	5	SRI AKASH U G	8000
5	6	SRI HARSHAVARDHANAN R	3000
6	7	SRI HARSHAVARDHANAN R	6000

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334) Experiment: 04

```
#sample calculation for low range(lr) , upper range (ur),percentile import  
numpy as np  
array=np.random.randint(1,100,16) # randomly generate 16 numbers between 1 to
```

```

100 array array([27, 50, 44, 6, 58, 61, 23, 86, 67, 20, 75, 7, 79, 61,
90, 54]) array.mean()
50.5 np.percentile(array,25)
26.0 np.percentile(array,50)
56.0 np.percentile(array,75)
69.0 np.percentile(array,100)
90.0

```

```


#outliers detection def
outDetection(array):
sorted(array)
Q1,Q3=np.percentile(array,[25,75])
IQR=Q3-Q1 lr=Q1-(1.5*IQR)
ur=Q3+(1.5*IQR) return lr,ur
lr,ur=outDetection(array) lr,ur
(-38.5, 133.5)

```


```

import seaborn as sns %matplotlib
inline sns.displot(array)
<seaborn.axisgrid.FacetGrid at
0x78f3291c2710>

```




```
sns.distplot(array)
sns.distplot(array)
<Axes: ylabel='Density'>
```




```
new_array=array[(array>lr) & (array<ur)] new_array array([27, 50, 44,
6, 58, 61, 23, 86, 67, 20, 75, 7, 79, 61, 90, 54])
```

```
sns.distplot(new_array)
<seaborn.axisgrid.FacetGrid at
0x78f2e09bb580>
```



```
lr1,ur1=outDetection(new_array)
lr1,ur1 (-38.5, 133.5)
final_array=new_array[(new_array>lr1) & (new_array<ur1)]
```

```
final_array array([27, 50, 44, 6, 58, 61, 23, 86, 67, 20, 75, 7, 79,
61, 90, 54]) sns.distplot(final_array)
```



Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 05

```
import numpy as np import pandas
```

```
as pd
```

```
df=pd.read_csv("Hotel_Dataset.csv"
```

```
)
```

```
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	lbys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
9	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

```
df.duplicated()
```

```

0    False
1    False
2    False
3    False
4    False
5    False
6    False
7    False
8    False
9     True
10   False
dtype: bool

```

df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11 entries, 0 to 10
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  --
0   CustomerID            11 non-null    int64
1   Age_Group             11 non-null    object
2   Rating(1-5)           11 non-null    int64
3   Hotel                 11 non-null    object
4   FoodPreference         11 non-null    object
5   Bill                  11 non-null    int64
6   NoOfPax               11 non-null    int64
7   EstimatedSalary       11 non-null    int64
8   Age_Group.1           11 non-null    object
dtypes: int64(5), object(4)
memory usage: 924.0+ bytes

```

df.drop_duplicates(inplace=True)

df len(df) 10

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

index=np.array(list(range(0,len(df))))

) df.set_index(index,inplace=True)

index

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

df

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
9	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

```
df.drop(['Age_Group.1'],axis=1,inplace=True)
```

df

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1	20-25	4	Ibis	veg	1300	2	40000
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000
2	3	25-30	6	RedFox	Veg	1322	2	30000
3	4	20-25	-1	LemonTree	Veg	1234	2	120000
4	5	35+	3	Ibis	Vegetarian	989	2	45000
5	6	35+	3	Ibys	Non-Veg	1909	2	122220
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122
7	8	20-25	7	LemonTree	Veg	2999	-10	345673
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999
9	10	30-35	5	RedFox	non-Veg	-6755	4	87777

```
df.CustomerID.loc[df.CustomerID<0]=np.nan
```

```
df.Bill.loc[df.Bill<0]=np.nan
```

```
df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan
```

df

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4.0	Ibis	veg	1300.0	2	40000.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0	3	59000.0
2	3.0	25-30	NaN	RedFox	Veg	1322.0	2	30000.0
3	4.0	20-25	NaN	LemonTree	Veg	1234.0	2	120000.0
4	5.0	35+	3.0	Ibis	Vegetarian	989.0	2	45000.0
5	6.0	35+	3.0	Ibys	Non-Veg	1909.0	2	122220.0
6	7.0	35+	4.0	RedFox	Vegetarian	1000.0	-1	21122.0
7	8.0	20-25	NaN	LemonTree	Veg	2999.0	-10	345673.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0	3	NaN
9	10.0	30-35	5.0	RedFox	non-Veg	NaN	4	87777.0

```
df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4.0	Ibis	veg	1300.0	2.0	40000.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0	3.0	59000.0
2	3.0	25-30	NaN	RedFox	Veg	1322.0	2.0	30000.0
3	4.0	20-25	NaN	LemonTree	Veg	1234.0	2.0	120000.0
4	5.0	35+	3.0	Ibis	Vegetarian	989.0	2.0	45000.0
5	6.0	35+	3.0	Ibys	Non-Veg	1909.0	2.0	122220.0
6	7.0	35+	4.0	RedFox	Vegetarian	1000.0	NaN	21122.0
7	8.0	20-25	NaN	LemonTree	Veg	2999.0	NaN	345673.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0	3.0	NaN
9	10.0	30-35	5.0	RedFox	non-Veg	NaN	4.0	87777.0

```
df.Age_Group.unique()
```

```
array(['20-25', '30-35', '25-30', '35+'], dtype=object)
```

```
df.Hotel.unique()
```

```
array(['Ibis', 'LemonTree', 'RedFox', 'Ibys'], dtype=object)
```

```
df.Hotel.replace(['Ibys'],'Ibis',inplace=True) df.FoodPreference.unique
```

```
<bound method Series.unique of 0 veg
```

```
1 Non-Veg
```

```
2 Veg
```

```
3 Veg
```

```
4 Vegetarian
```

```
5 Non-Veg
```

```
6 Vegetarian
```

```

7 Veg
8 Non-Veg
9 non-Veg
Name: FoodPreference, dtype: object>
df.FoodPreference.replace(['Vegetarian','veg'],'Veg',inplace=True)

df.FoodPreference.replace(['non-Veg'],'Non-Veg',inplace=True)

df.EstimatedSalary.fillna(round(df.EstimatedSalary.mean()),inplace=True)

df.NoOfPax.fillna(round(df.NoOfPax.median()),inplace=True)

df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()), inplace=True)

df.Bill.fillna(round(df.Bill.mean()),inplace=True) df

```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4.0	Ibis	Veg	1300.0	2.0	40000.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0	3.0	59000.0
2	3.0	25-30	4.0	RedFox	Veg	1322.0	2.0	30000.0
3	4.0	20-25	4.0	LemonTree	Veg	1234.0	2.0	120000.0
4	5.0	35+	3.0	Ibis	Veg	989.0	2.0	45000.0
5	6.0	35+	3.0	Ibis	Non-Veg	1909.0	2.0	122220.0
6	7.0	35+	4.0	RedFox	Veg	1000.0	2.0	21122.0
7	8.0	20-25	4.0	LemonTree	Veg	2999.0	2.0	345673.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0	3.0	96755.0
9	10.0	30-35	5.0	RedFox	Non-Veg	1801.0	4.0	87777.0

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 06

```

import numpy as np import
pandas as pd
df=pd.read_csv('/content/pre-process_datasample.csv')

```

df

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	NaN	Yes
5	France	35.0	58000.0	Yes
6	Spain	NaN	52000.0	No
7	France	48.0	79000.0	Yes
8	NaN	50.0	83000.0	No
9	France	37.0	67000.0	Yes

Next steps: df.head()

Country Age Salary Purchased 0

France 44.0 72000.0 No **1** Spain 27.0
48000.0 Yes **2** Germany 30.0 54000.0
No **3** Spain 38.0 61000.0 No **4**
Germany 40 0 NaN Yes

```
df.Country.fillna(df.Country.mode()[0],inplace=True) features=df.iloc[:,:-1].values
```

```
df.Country.fillna(df.Country.mode()[0],inplace=True)
```

```
label=df.iloc[:,-1].values
```

```

from sklearn.impute import SimpleImputer

age=SimpleImputer(strategy="mean",missing_values=np.nan)

Salary=SimpleImputer(strategy="mean",missing_values=np.nan)

age.fit(features[:,[1]])

```

```

    ▼ SimpleImputer
SimpleImputer()

```

```

Salary.fit(features[:,[2]])

    ▼ SimpleImputer SimpleImputer()

```

```

SimpleImputer()

    ▼ SimpleImputer SimpleImputer()
features[:,[1]]=age.transform(features[:,[1]])

```

```

features[:,[2]]=Salary.transform(features[:,[2]])

```

```

features

array([[ 'France', 44.0, 72000.0],
 ['Spain', 27.0, 48000.0],
 ['Germany', 30.0, 54000.0],
 ['Spain', 38.0, 61000.0],
 ['Germany', 40.0, 63777.77777777778],
 ['France', 35.0, 58000.0],
 ['Spain', 38.77777777777778, 52000.0],
 ['France', 48.0, 79000.0],
 ['France', 50.0, 83000.0],
 ['France', 37.0, 67000.0]], dtype=object)

```

```

from sklearn.preprocessing import OneHotEncoder

oh = OneHotEncoder(sparse_output=False)

Country=oh.fit_transform(features[:,[0]])

```

Country

```
array([[1., 0., 0.],
       [0., 0., 1.],
       [0., 1., 0.],
       [0., 0., 1.],
       [0., 1., 0.],
       [1., 0., 0.],
       [0., 0., 1.],
       [1., 0., 0.],

       [1., 0., 0.],
       [1., 0., 0.]])
```

```
final_set=np.concatenate((Country,features[:,[1,2]]),axis=1)
```

final_set

```
array([[1.0, 0.0, 0.0, 44.0, 72000.0],
       [0.0, 0.0, 1.0, 27.0, 48000.0],
       [0.0, 1.0, 0.0, 30.0, 54000.0],
       [0.0, 0.0, 1.0, 38.0, 61000.0],
       [0.0, 1.0, 0.0, 40.0, 63777.777777777778],
       [1.0, 0.0, 0.0, 35.0, 58000.0],
       [0.0, 0.0, 1.0, 38.77777777777778, 52000.0],
       [1.0, 0.0, 0.0, 48.0, 79000.0],
       [1.0, 0.0, 0.0, 50.0, 83000.0],
       [1.0, 0.0, 0.0, 37.0, 67000.0]], dtype=object)
```

```
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
sc.fit(final_set)
feat_standard_scaler=sc.transform(final_set)
```

feat_standard_scaler

```
array([[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
         7.58874362e-01,  7.49473254e-01],
       [-1.00000000e+00, -5.00000000e-01,  1.52752523e+00,
        -1.71150388e+00, -1.43817841e+00],
       [-1.00000000e+00,  2.00000000e+00, -6.54653671e-01,
        -1.27555478e+00, -8.91265492e-01],
       [-1.00000000e+00, -5.00000000e-01,  1.52752523e+00,
        -1.13023841e-01, -2.53200424e-01],
       [-1.00000000e+00,  2.00000000e+00, -6.54653671e-01,
```

```

1.77608893e-01, 6.63219199e-16],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
-5.48972942e-01, -5.26656882e-01],
[-1.00000000e+00, -5.00000000e-01, 1.52752523e+00,
0.00000000e+00, -1.07356980e+00],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
1.34013983e+00, 1.38753832e+00],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
1.63077256e+00, 1.75214693e+00],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
-2.58340208e-01, 2.93712492e-01]])
from sklearn.preprocessing import MinMaxScaler
mms=MinMaxScaler(feature_range=(0,1))
mms.fit(final_set)
feat_minmax_scaler=mms.transform(final_set)
feat_minmax_scaler

array([[1. , 0. , 0. , 0.73913043, 0.68571429],
[0. , 0. , 1. , 0. , 0. ],
[0. , 1. , 0. , 0.13043478, 0.17142857],
[0. , 0. , 1. , 0.47826087, 0.37142857],
[0. , 1. , 0. , 0.56521739, 0.45079365],
[1. , 0. , 0. , 0.34782609, 0.28571429],
[0. , 0. , 1. , 0.51207729, 0.11428571],
[1. , 0. , 0. , 0.91304348, 0.88571429],
[1. , 0. , 0. , 1. , 1. ],
[1. , 0. , 0. , 0.43478261, 0.54285714]])

```

Lab experiments

Roll no: 230701038

Name:Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 07

```

import numpy as np
import pandas as pd
df=pd.read_csv("/content/pre-process_datasample.csv") df

```

Country Age Salary Purchased

0 France 44.0 72000.0 No

1 Spain 27.0 48000.0 Yes

2 Germany 30.0 54000.0 No

3 Spain 38.0 61000.0 No
4 Germany 40.0 NaN Yes
5 France 35.0 58000.0 Yes **6** Spain NaN 52000.0 No
7 France 48.0 79000.0 Yes **8**
 NaN 50.0 83000.0 No
9 France 37.0 67000.0 Yes

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9  Data
columns (total 4 columns):
# Column Non-Null Count Dtype
---  ---
0    Country 9 non-null object
1    Age 9 non-null float64
2    Salary 9 non-null float64  3 Purchased 10
non-null object  dtypes: float64(2), object(2)
memory usage: 448.0+ bytes  df.Country.mode()

```

```

Country 0
France

```

```
df.Country.mode()[0]
```

```
type(df.Country.mode())
```

```
df.Country.fillna(df.Country.mode()[0],inplace=True)
```

```
df.Age.fillna(df.Age.median(),inplace=True)
```

```
df.Salary.fillna(round(df.Salary.mean()),inplace=True)
```

```
df
```

```

Country Age Salary Purchased
0 France 44.0 72000.0 No
1 Spain 27.0 48000.0 Yes

```

2 Germany 30.0 54000.0 No
3 Spain 38.0 61000.0 No
4 Germany 40.0 63778.0 Yes
5 France 35.0 58000.0 Yes
6 Spain 38.0 52000.0 No
7 France 48.0 79000.0 Yes **8** France 50.0 83000.0 No
9 France 37 0 67000 0 Yes

```
pd.get_dummies(df.Country)
```

France Germany Spain
0 True False False
1 False False True
2 False True False
3 False False True
4 False True False
5 True False False
6 False False True
7 True False False
8 True False False
9 True False False

```
updated_dataset=pd.concat([pd.get_dummies(df.Country),df.iloc[:,[1,2,3]]],axis=1)
```

```
updated_dataset
```

France Germany Spain Age Salary Purchased
0 True False False 44.0 72000.0 No
1 False False True 27.0 48000.0 Yes
2 False True False 30.0 54000.0 No
3 False False True 38.0 61000.0 No
4 False True False 40.0 63778.0 Yes
5 True False False 35.0 58000.0 Yes
6 False False True 38.0 52000.0 No
7 True False False 48.0 79000.0 Yes **8** True False False 50.0 83000.0 No
9 True False False 37 0 67000 0 Yes

```
df.info()
```



```
updated_dataset.Purchased.replace(['No', 'Yes'], [0,1], inplace=True)
```

```
updated_dataset
```

```
   France Germany Spain Age  Salary  Purchased
0  True  False  False  44.0  72000.0  0
1  False  False   True  27.0  48000.0  1
2  False   True  False  30.0  54000.0  0
3  False  False   True  38.0  61000.0  0
4  False   True  False  40.0  63778.0  1
5  True  False  False  35.0  58000.0  1
6  False  False   True  38.0  52000.0  0
7  True  False  False  48.0  79000.0  1
8  True  False  False  50.0  83000.0  0
9  True  False  False  37.0  67000.0  1
```

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

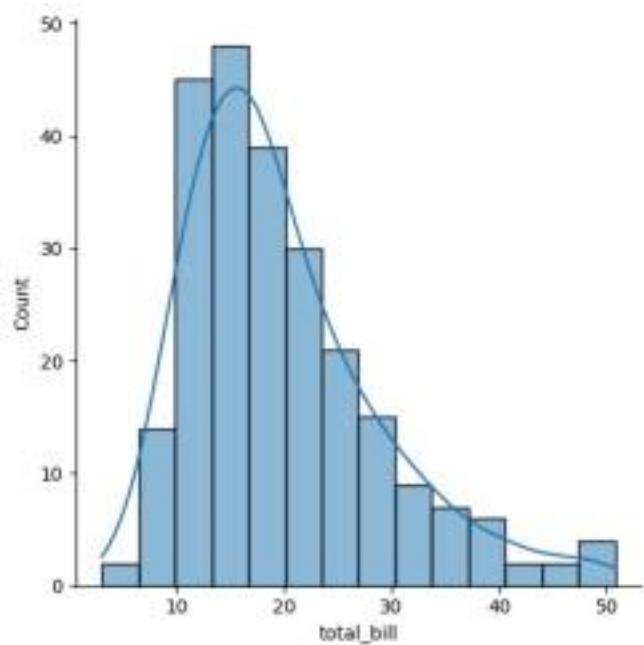
Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

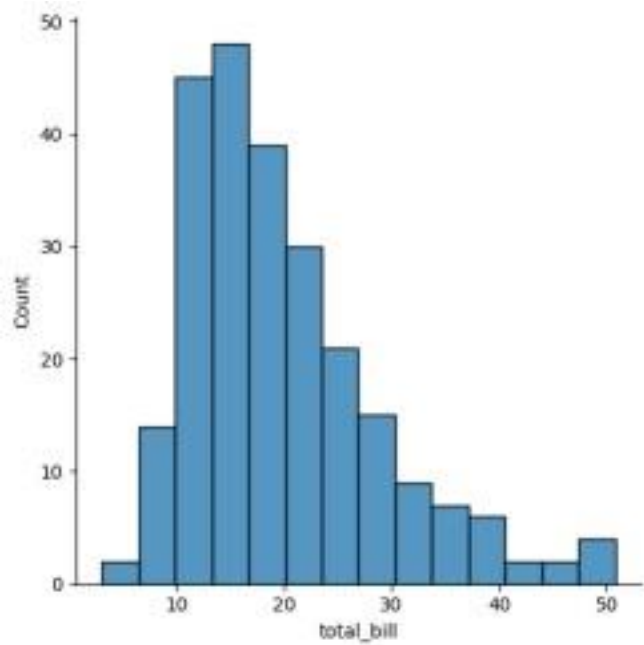
Experiment: 08

```
import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
tips=sns.load_dataset('tips')
tips.head()
```

```
   total_bill  tip  sex  smoker  day  time  size
0    16.99  1.01 Female    No  Sun  Dinner    2
1    10.34  1.66  Male    No  Sun  Dinner    3
2    21.01  3.50  Male    No  Sun  Dinner    3
3    23.68  3.31  Male    No  Sun  Dinner    2
4    24.59  3.61 Female    No  Sun  Dinner    4
sns.displot(tips.total_bill, kde=True)
<seaborn.axisgrid.FacetGrid at 0x79bb4c7ea680>
```




```
sns.displot(tips.total_bill,kde=False)
```




```
sns.jointplot(x=tips.tip,y=tips.total_bill)
```


```
<seaborn.axisgrid.JointGrid at 0x79bb08fc96c0>
```



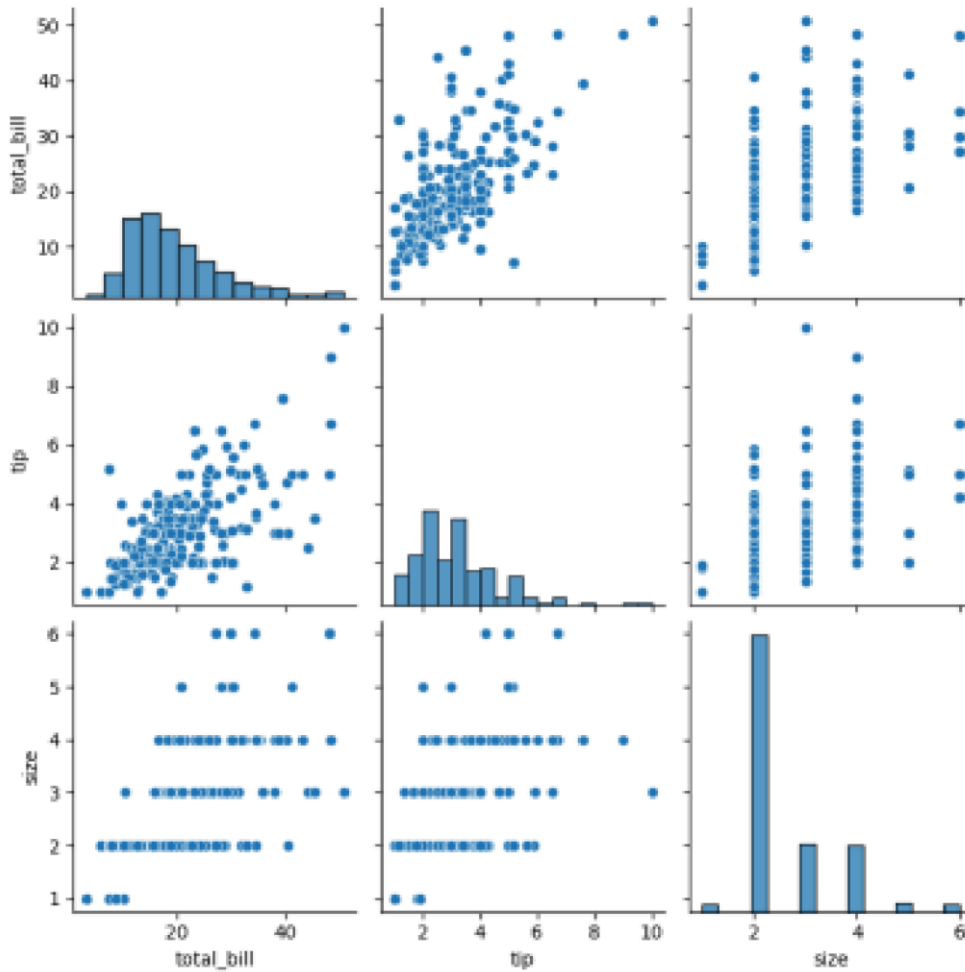
```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="reg")
```



```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="hex")  
<seaborn.axisgrid.JointGrid at 0x79bb088f4730>
```



```
sns.pairplot(tips)
```



```
tips.time.value_counts()
```

```
count
```

```
time
```

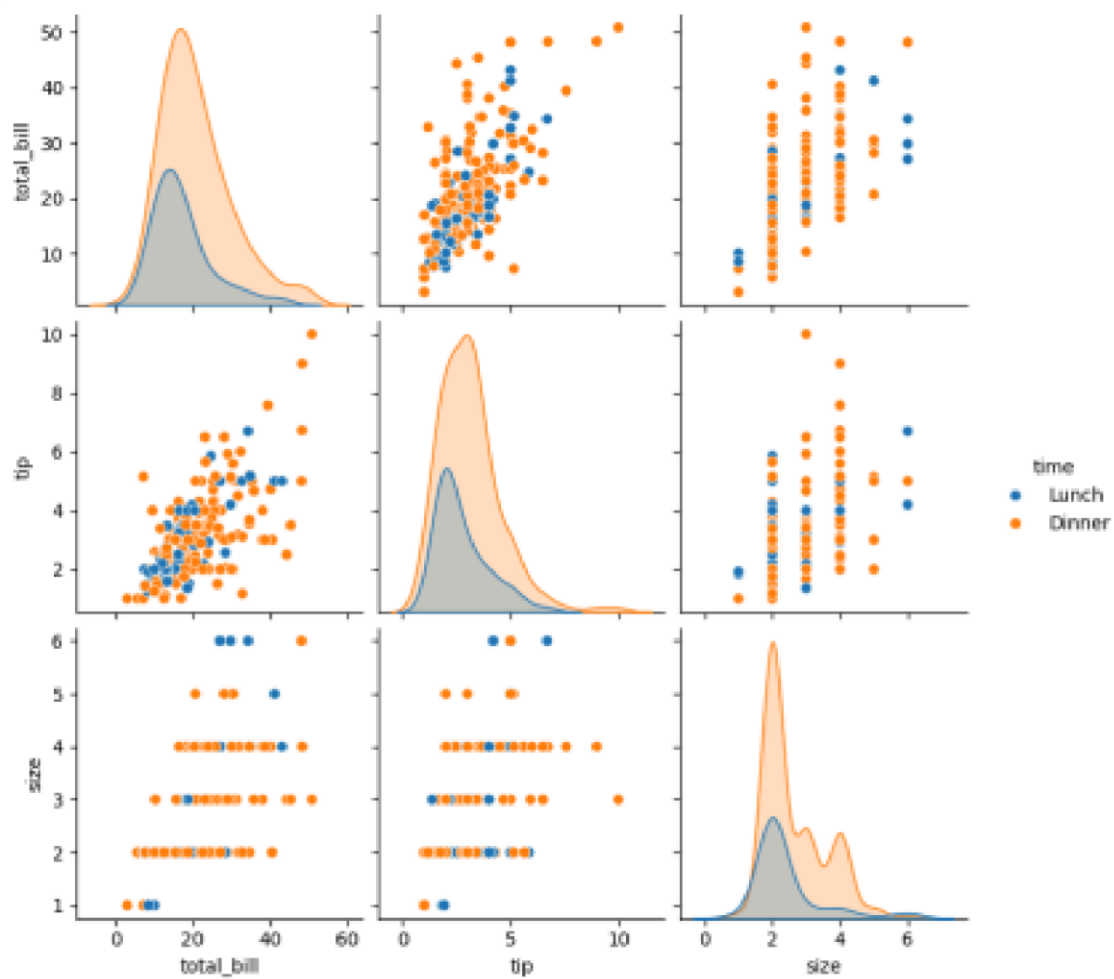
```
Dinner 176
```

```
Lunch 68 dtype:
```

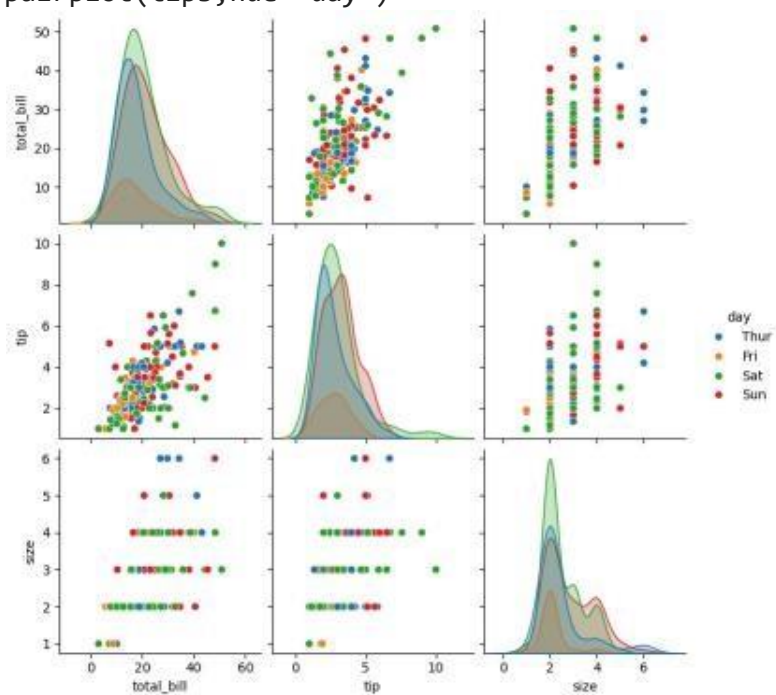
```
int64
```

```
sns.pairplot(tips,hue='time')
```

```
<seaborn.axisgrid.PairGrid at 0x79bb088f4670>
```




```
sns.pairplot(tips, hue='day')
```



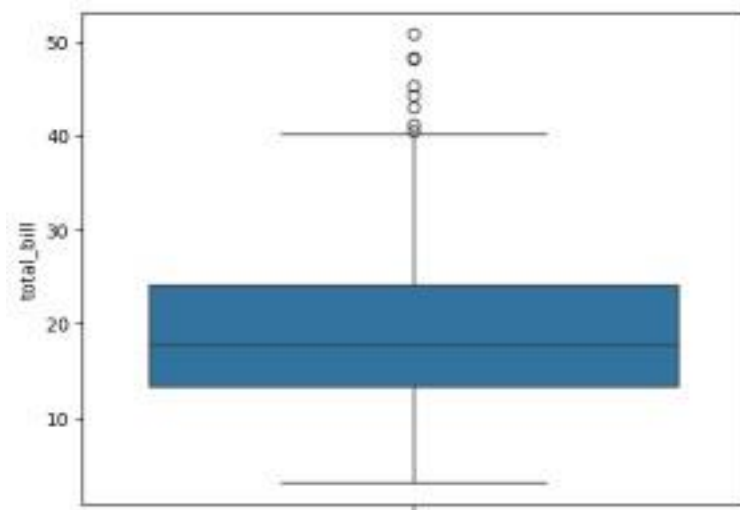
```
sns.heatmap(tips.corr(numeric_only=True),annot=True)
```

<Axes: >



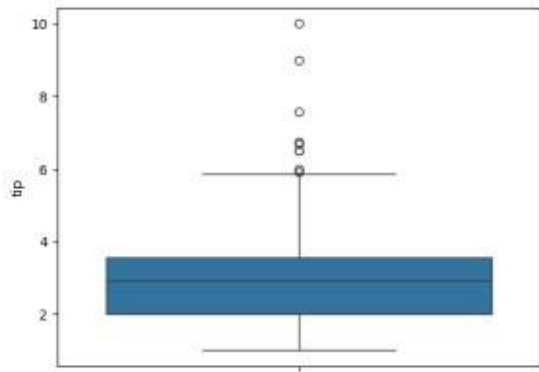
```
sns.boxplot(tips.total_bill)
```

<Axes: ylabel='total_bill'>




```
sns.boxplot(tips.tip)
```

<Axes: ylabel='tip'>



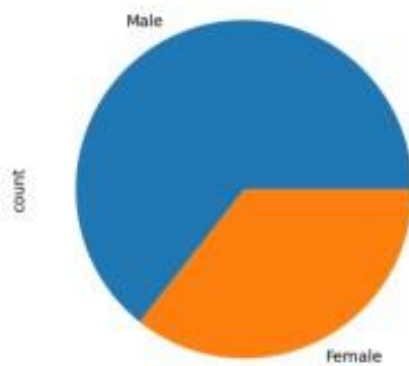

```
sns.countplot(tips.day)
```

```
<Axes: xlabel='count', ylabel='day'>
```




```
sns.countplot(tips.sex)
```

```
h<Axes: xlabel='count', ylabel='sex'>
```

```
tips.sex.value_counts().plot(kind='pie')  
<Axes: ylabel='count'>
```

```
tips.sex.value_counts().plot(kind='bar')  
<Axes: xlabel='sex'>
```



Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 09

```
df.dropna(inplace=True) df.info()
```

```
<class 'pandas.core.frame.DataFrame'> RangeIndex: 30 entries, 0 to 29
```

```
Data columns (total 2 columns):
```

```
df.describe()
```

```
Out[5]:
```

	Years	Experience	Salary	count
	30.000000	30.000000	30.000000	30.000000
	mean	5.313333	76003.000000	std
	2.837888	27414.429785		
	min	1.100000	37731.000000	
	25%	3.200000	56720.750000	
	50%	4.700000	65237.000000	
	75%	7.700000	100544.750000	
	max	10.500000	122391.000000	

```
In [6]: features=df.iloc[:,[0]].values
```

```
label=df.iloc[:,[1]].values
```

```

from sklearn.model_selection
import train_test_split
x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,random_
st

from sklearn.linear_model

import LinearRegression model=LinearRegression()
model.fit(x_train,y_train)

Out[20]: ▼ LinearRegression
          LinearRegression()

          model.score(x_tr
In [21]:      ain,y_train)
Out[21]: 0.9603182547438908
          model.score(x_t
In [23]:      est,y_test)
Out[23]: 0.9184170849214232
          model.coef_ In [24]:      f_
Out[24]:
          array([[9281.30847068]])
          model.intercept_ In [25]:      cept_
Out[25]: array([27166.73682891])

In [26]:
import pickle
pickle.dump(model,open('SalaryPred.model','wb'))

model=pickle.load(open('SalaryPred.model','rb')) yr_of_exp=float(input("Enter Years
of Experience: "))

yr_of_exp_NP=np.array([[yr_of_exp]])
Salary=model.predict(yr_of_exp_NP)
Enter Years of Experience: 44

print("Estimated Salary for {} years of experience is {}:"

```

```
.format(yr_of_exp,Salary) Estimated Salary for 44.0 years of experience is
```

```
[[435544.30953887]]:
```

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 10

```
import numpy as np
```

```
import pandas as pd
```

```
df=pd.read_csv('Iris.csv') df.info()
```

```
df.variety.value_counts()
```

```
Out[3]: Setosa 50  
        Versicolor 50  
        Virginica 50  
        Name: variety, dtype: int64
```

```
In [4]: df.head()
```

```
Out[4]: sepal.length sepal.width petal.length petal.width variety 0 5.1 3.5 1.4 0.2 Setosa  
        1 4.9 3.0 1.4 0.2 Setosa 2 4.7 3.2 1.3 0.2 Setosa 3 4.6 3.1 1.5  
        0.2 Setosa 4 5.0 3.6 1.4 0.2 Setosa
```

```
In [5]: In [6]: In [8]:
```

```
features=df.iloc[:, :-1].values label=df.iloc[:, 4].values
```

```
from sklearn.model_selection
```

```
import train_test_split from sklearn.neighbors
```

```
import KNeighborsClassifier
```

```
xtrain,xtest,ytrain,ytest=train_test_split(features,label,test_size=.2,random_state=42)  
model_KNN=KNeighborsClassifier(n_neighbors=5)
```

```
model_KNN.fit(xtrain,ytrain)
```

```
Out[8]: KNeighborsClassifier()
print(model_KNN.score(xtrain,ytrain))
print(model_KNN.score(xtest,ytest))

0.9583333333333334
1.0
```

```
from sklearn.metrics import confusion_matrix
confusion_matrix(label,model_KNN.predict(features))
```

```
Out[10]: array([[50, 0, 0],
               [ 0, 47, 3],
               [ 0, 2, 48]], dtype=int64) from sklearn.metrics import
classification_report
print(classification_report(label,model_KNN.predict(features))
) precision recall f1-score support

Setosa 1.00 1.00 1.00 50 Versicolor 0.96 0.94 0.95 50 Virginica
0.94 0.96 0.95 50

accuracy 0.97 150 macro avg 0.97 0.97 0.97 150 weighted avg 0.97
0.97 0.97 150
```

Lab experiments
Roll no: 2307010138
Name: Arvind Ravi Class: CSE-A III
Subject: Fundamentals of data science (CS2334)
Experiment: 11

```
In [1]: import numpy as np import pandas
as pd
df=pd.read_csv('Social_Network_Ads.csv') df
```

```
Out[1]: User ID Gender Age EstimatedSalary Purchased 0 15624510 Male 19 19000 0
1 15810944 Male 35 20000 0 2 15668575 Female 26 43000
0 3 15603246 Female 27 57000 0 4 15804002 Male 19
76000 0 ... ..
395 15691863 Female 46 41000 1 396 15706071 Male 51 23000 1
397 15654296 Female 50 20000 1 398 15755018
Male 36 33000 0 399 15594041 Female 49 36000 1

400 rows × 5 columns
```

```
In [2]: df.head()
```

Out[2]: User ID Gender Age EstimatedSalary Purchased

0 15624510 Male 19 19000 0
1 15810944 Male 35 20000 0
2 15668575 Female 26 43000 0
3 15603246 Female 27 57000 0
4 15804002 Male 19 76000 0

```
In [4]: features=df.iloc[:,[2,3]].values
label=df.iloc[:,4].values features
```

```
Out[4]: array([[ 19, 19000], [ 35,
20000],
[ 26, 43000],
[ 27, 57000],
[ 19, 76000],
[ 27, 58000],
[ 27, 84000],
[ 32, 150000],
[ 25, 33000],
[ 35, 65000],
[ 26, 80000],
[ 26, 52000],
[ 20, 86000],
[ 32, 18000],
[ 18, 82000],
[ 29, 80000],
[ 47, 25000],
[ 45, 26000],
[ 46, 28000],
[ 48 29000])
```

```
In [5]: label
```

[illegible]

```

1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0,
1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1,
0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1,
0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1,
0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1,
1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1], dtype=int64)

```

```

In [6]: from sklearn.model_selection import
train_test_split
from sklearn.linear_model import LogisticRegression

```

```

for i in range(1,401):

```

```

x_train,x_test,y_train,y_test=train_test_split(features,label,
test_size=0. model=LogisticRegression())
model.fit(x_train,y_train)
train_score=model.score(x_train,y_train)
test_score=model.score(x_test,y_test) if
test_score>train_score:
    print("Test {} Train{} Random State
{}".format(test_score,train_score,i)

```

```

Test 0.6875 Train0.63125 Random State 3
Test 0.7375 Train0.61875 Random State 4
Test 0.6625 Train0.6375 Random State 5
Test 0.65 Train0.640625 Random State 6
Test 0.675 Train0.634375 Random State 7
Test 0.675 Train0.634375 Random State 8
Test 0.65 Train0.640625 Random State 10
Test 0.6625 Train0.6375 Random State 11
Test 0.7125 Train0.625 Random State 13
Test 0.675 Train0.634375 Random State 16
Test 0.7 Train0.628125 Random State 17
Test 0.7 Train0.628125 Random State 21
Test 0.65 Train0.640625 Random State 24
Test 0.6625 Train0.6375 Random State 25
Test 0.75 Train0.615625 Random State 26
Test 0.675 Train0.634375 Random State 27 Test
0.7 Train0.628125 Random State 28
Test 0.6875 Train0.63125 Random State 29
Test 0.6875 Train0.63125 Random State 31 T
t 0 6625 T i 0 6375 R d St t 37

```

```

x_train,x_test,y_train,y_test=train_test_split(features,label,
test_size=0.2, finalModel=LogisticRegression())
finalModel.fit(x_train,y_train)

```

```
Out[8]: LogisticRegression()
```

```
print(finalModel.score(x_train,y_train))
```

```
print(finalModel.score(x_test,y_test))
```

```
0.834375
```

```
0.9125
```

```
from sklearn.metrics import classification_report
```

```
print(classification_report(label,finalModel.predict(features))
```

```
) precision recall f1-score support    0 0.85 0.93 0.89 257  1
```

```
0.84 0.71 0.77 143
```

```
accuracy 0.85 400  macro avg 0.85 0.82 0.83 400 weighted avg 0.85 0.85  
0.85 400
```

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 12

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
%matplotlib inline
```

```
df=pd.read_csv('Mall_Customers  
.csv') df.info()
```

```
<class 'pandas.core.frame.DataFrame'> RangeIndex:
```

```
200 entries, 0 to 199
```

```
Data columns (total 5 columns):
```

```
df.head()
```

```
Out[4]: CustomerID Gender Age Annual Income (k$) Spending Score (1-100)
```

```
0 1 Male 19 15 39
```

```
1 2 Male 21 15 81
```

```
2 3 Female 20 16 6
```

```
3 4 Female 23 16 77
```


4 5 Female 31 17 40

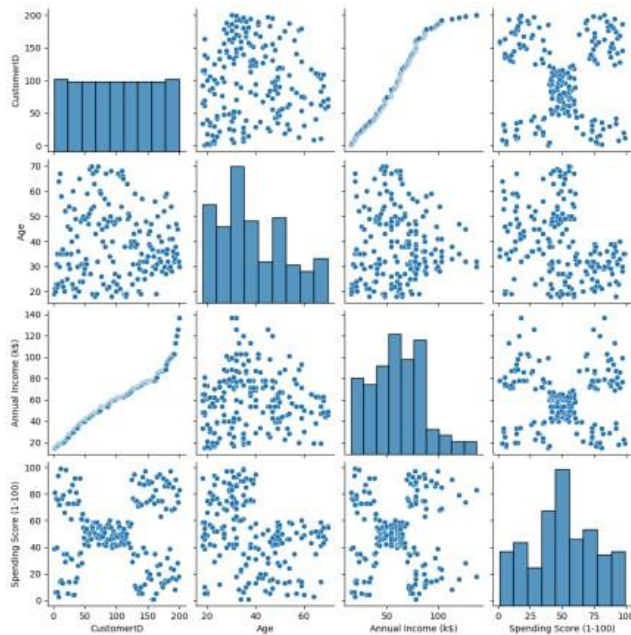
```
sns.pairplot(df)
```

In [5]:

Out[5]: <seaborn.axisgrid.PairGrid at 0x170e8e47850>

In [6]:

```
features=df.iloc[:,[3,4]].values
```



```
In [7]: from sklearn.cluster
import KMeans
model=KMeans(n_clusters=5)
model.fit(features)
KMeans(n_clusters=5)
```

Out[7]: KMeans(n_clusters=5)

```
In [8]:
Final=df.iloc[:,[3,4]]
Final['label']=model.predict(features)
```

```
Final.head()
Final['label']=model.predict(features)
```

Out[8]: Annual Income (k\$) Spending Score (1-100) label


0 15 39 4

1 15 81 2 2 16 6 4

3 16 77 2

4 17 40 4


```
In [9]: sns.set_style("whitegrid") sns.FacetGrid(Final,hue="label",height=8)
\
.map(plt.scatter,"Annual Income (k$)", "Spending Score (1-100)") \
.add_legend();
plt.show()
```



In [10]:

```
features_el=df.iloc[:,[2,3,4]].values
from sklearn.cluster import KMeans
wcss=[] for i in range(1,10):
    model=KMeans(n_clusters=i)
    model.fit(features_el)
    wcss.append(model.inertia_)
plt.plot(range(1,10),wcss)
```

Out[10]: [<matplotlib.lines.Line2D at 0x170e99f3550>]



Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 13

```
import numpy as np
import matplotlib.pyplot as plt
```

```
# Step 1: Generate a population (e.g., normal distribution) population_mean =
50
population_std = 10
```

```

population_size = 100000
population = np.random.normal(population_mean, population_std,
population_size)
# Step 2: Random sampling
sample_sizes = [30, 50, 100]
num_samples = 1000
sample_means = {}
for size in sample_sizes:
    sample_means[size] = []

    for _ in range(num_samples):
        sample = np.random.choice(population, size=size, replace=False)


        sample_means[size].append(np.mean(sample))

plt.figure(figsize=(12, 8))
for i, size in enumerate(sample_sizes):
    plt.subplot(len(sample_sizes), 1, i+1)
    plt.hist(sample_means[size], bins=30, alpha=0.7, label=f'Sample Size {size}')
    plt.axvline(np.mean(population), color='red', linestyle='dashed', linewidth=1.5, label='Population Mean')
    plt.title(f'Sampling Distribution (Sample Size {size})')
    plt.xlabel('Sample Mean')
    plt.ylabel('Frequency')
    plt.legend()

plt.tight_layout()
plt.show()

```

OUTPUT:



Lab experiments

Roll no: 230701038
 Name: Arvind Ravi
 Class: CSE-A III
 Subject: Fundamentals of data science (CS2334)
 Experiment: 13

```
import numpy as np
import scipy.stats as stats

sample_data = np.array([152, 148, 151, 149, 147, 153, 150, 148, 152,
                        149, 151, 150, 149, 152, 151, 148, 150, 152,
                        149, 150, 148, 153, 151, 150, 149, 152,
                        148, 151, 150, 153])

population_mean = 150

sample_mean = np.mean(sample_data)

sample_std = np.std(sample_data, ddof=1)

n = len(sample_data)

z_statistic = (sample_mean - population_mean) / (sample_std / np.sqrt(n))

p_value = 2 * (1 - stats.norm.cdf(np.abs(z_statistic)))

print(f"Sample Mean: {sample_mean:.2f}")

print(f"Z-Statistic: {z_statistic:.4f}")
```

```
print(f"P-Value:{p_value:.4f}")
```

```
alpha = 0.05 if p_value < alpha:    print("Reject the null hypothesis: The average weight is significantly different  
from 150 grams.") else:    print("Fail to reject the null hypothesis: There is no significant difference in average  
weight from 150 grams.")
```

OUTPUT:

Sample Mean: 150.20

Z-Statistic: 0.6406

P-Value: 0.5218

Fail to reject the null hypothesis: There is no significant difference in average weight from 150 grams.

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 14

```
import numpy as np import  
scipy.stats as stats  
np.random.seed(42)
```

```
sample_size = 25 sample_data = np.random.normal(loc=102, scale=15,  
size=sample_size)
```

```
# Mean IQ of 102, SD of 15
```

```
hypothesis population_mean = 100
```

```
sample_mean = np.mean(sample_data) sample_std = np.std(sample_data, ddof=1)
```

```
n = len(sample_data)
```

```
p_value = stats.ttest_1samp(sample_data, population_mean)
```

```
print(f"Sample Mean: {sample_mean:.2f}")
print(f"T-Statistic: {t_statistic:.4f}")
print(f"P-Value: {p_value:.4f}")

if p_value < alpha:
    print("Reject the null hypothesis: The average IQ score is significantly
different from 100.")
else:
    print("Fail to reject the null hypothesis: There is no significant
difference in average IQ score from 100.")
```

OUTPUT:

```
Sample Mean: 99.56
T-Statistic: -0.1578
P-Value: 0.8761
Fail to reject the null hypothesis: There is no significant difference in average
IQ score from 100.
```

Lab experiments

Roll no: 230701038

Name: Arvind Ravi

Class: CSE-A III

Subject: Fundamentals of data science (CS2334)

Experiment: 15

```
import numpy as np
import scipy.stats as stats
```

```
np.random.seed(42)
```

```
(A, B, C) n_plants = 25
```

```

growth_A = np.random.normal(loc=10, scale=2, size=n_plants) growth_B
= np.random.normal(loc=12, scale=3, size=n_plants) growth_C =
np.random.normal(loc=15, scale=2.5, size=n_plants)

all_data = np.concatenate([growth_A, growth_B, growth_C])

treatment_labels = ['A'] * n_plants + ['B'] * n_plants + ['C'] * n_plants

p_value = stats.f_oneway(growth_A, growth_B, growth_C)

print("Treatment A Mean Growth:", np.mean(growth_A))
print("Treatment B Mean Growth:", np.mean(growth_B))
print("Treatment C Mean Growth:", np.mean(growth_C)) print()
print(f"F-Statistic: {f_statistic:.4f}") print(f"P-Value:
{p_value:.4f}")

if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference in
    mean growth rates among the three treatments.")
else:
    print("Fail to reject the null hypothesis: There is no significant
    difference in mean growth rates among the three treatments.")

if p_value < alpha:      from statsmodels.stats.multicomp
import pairwise_tukeyhsd

    tukey_results = pairwise_tukeyhsd(all_data, treatment_labels, alpha=0.05)
print("\nTukey's HSD Post-hoc Test:")      print(tukey_results)

```

OUTPUT:

```

Treatment A Mean Growth: 9.672983882683818
Treatment B Mean Growth: 11.137680744437432
Treatment C Mean Growth: 15.265234904828972

```

F-Statistic: 36.1214

P-Value: 0.0000

Reject the null hypothesis: There is a significant difference in mean growth rates among the three treatments.